

Embedded realtime computer vision

**Victor Erukhimov
2016**

Emerging mobile/embedded use cases

Computational photography



Robotics



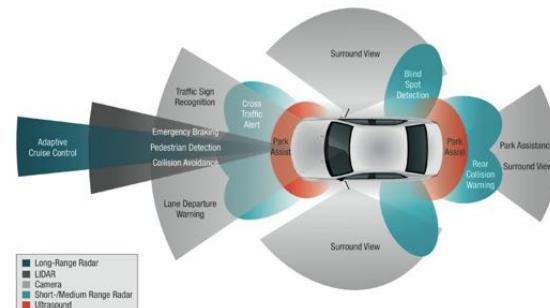
Drones



Wearable augmented reality



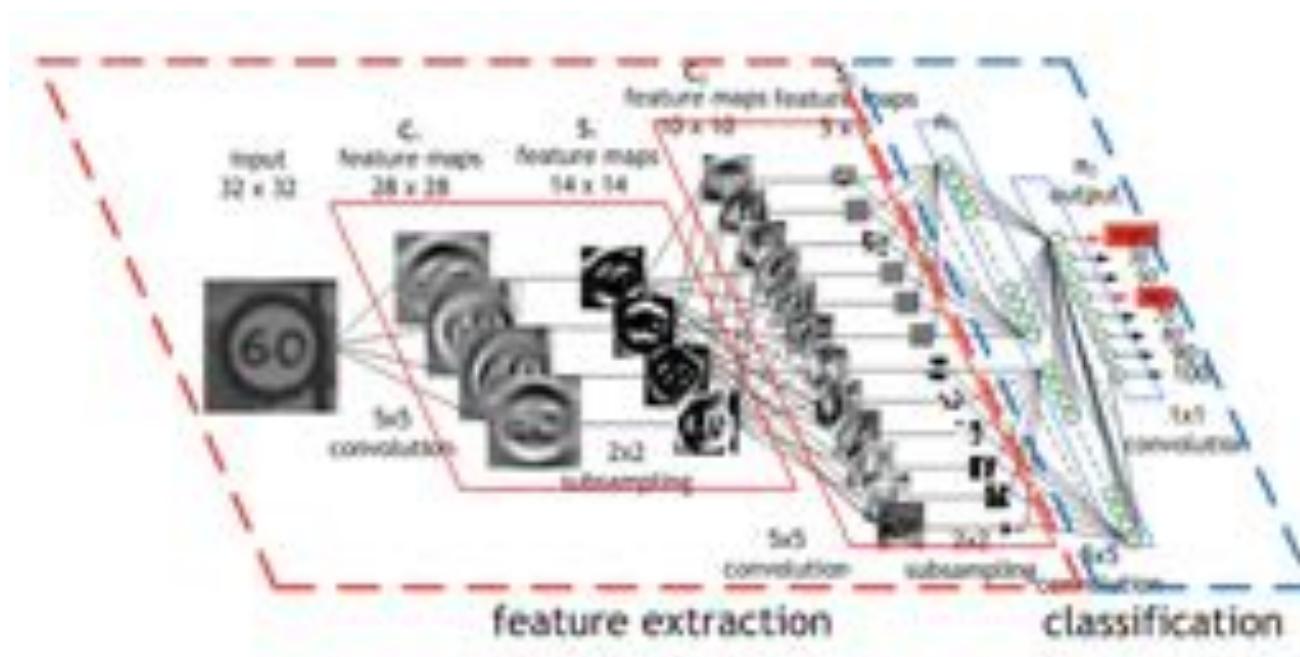
Driver assistance / driverless cars



Wearable virtual reality

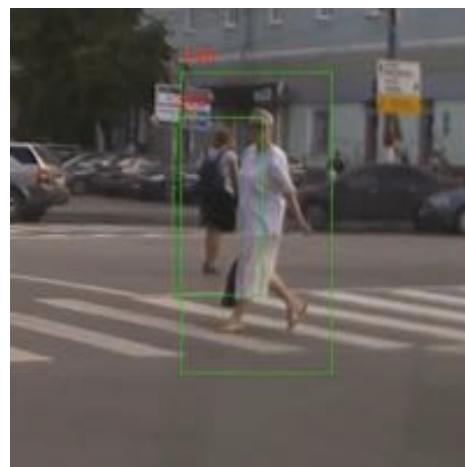


Convolutional neural networks



Detection involves a lot of convolution operations, that cannot be done in realtime on conventional embedded hardware

Example: driver assistance



Safety standards:

- MISRA C
- ISO 26262

Example: 3D scanning



Sample people models



ITSEEZ3D SCANNER. CONFIDENTIAL.

Sample object models



3D reconstruction on a tablet: full body



Reconstruction on a tablet, 45
seconds (iPad Air 2)



Cloud reconstruction, 10 minutes

Motivation

2010: OpenCV experimental port to
Android

Back camera:

- 5MP static images
- 480p video

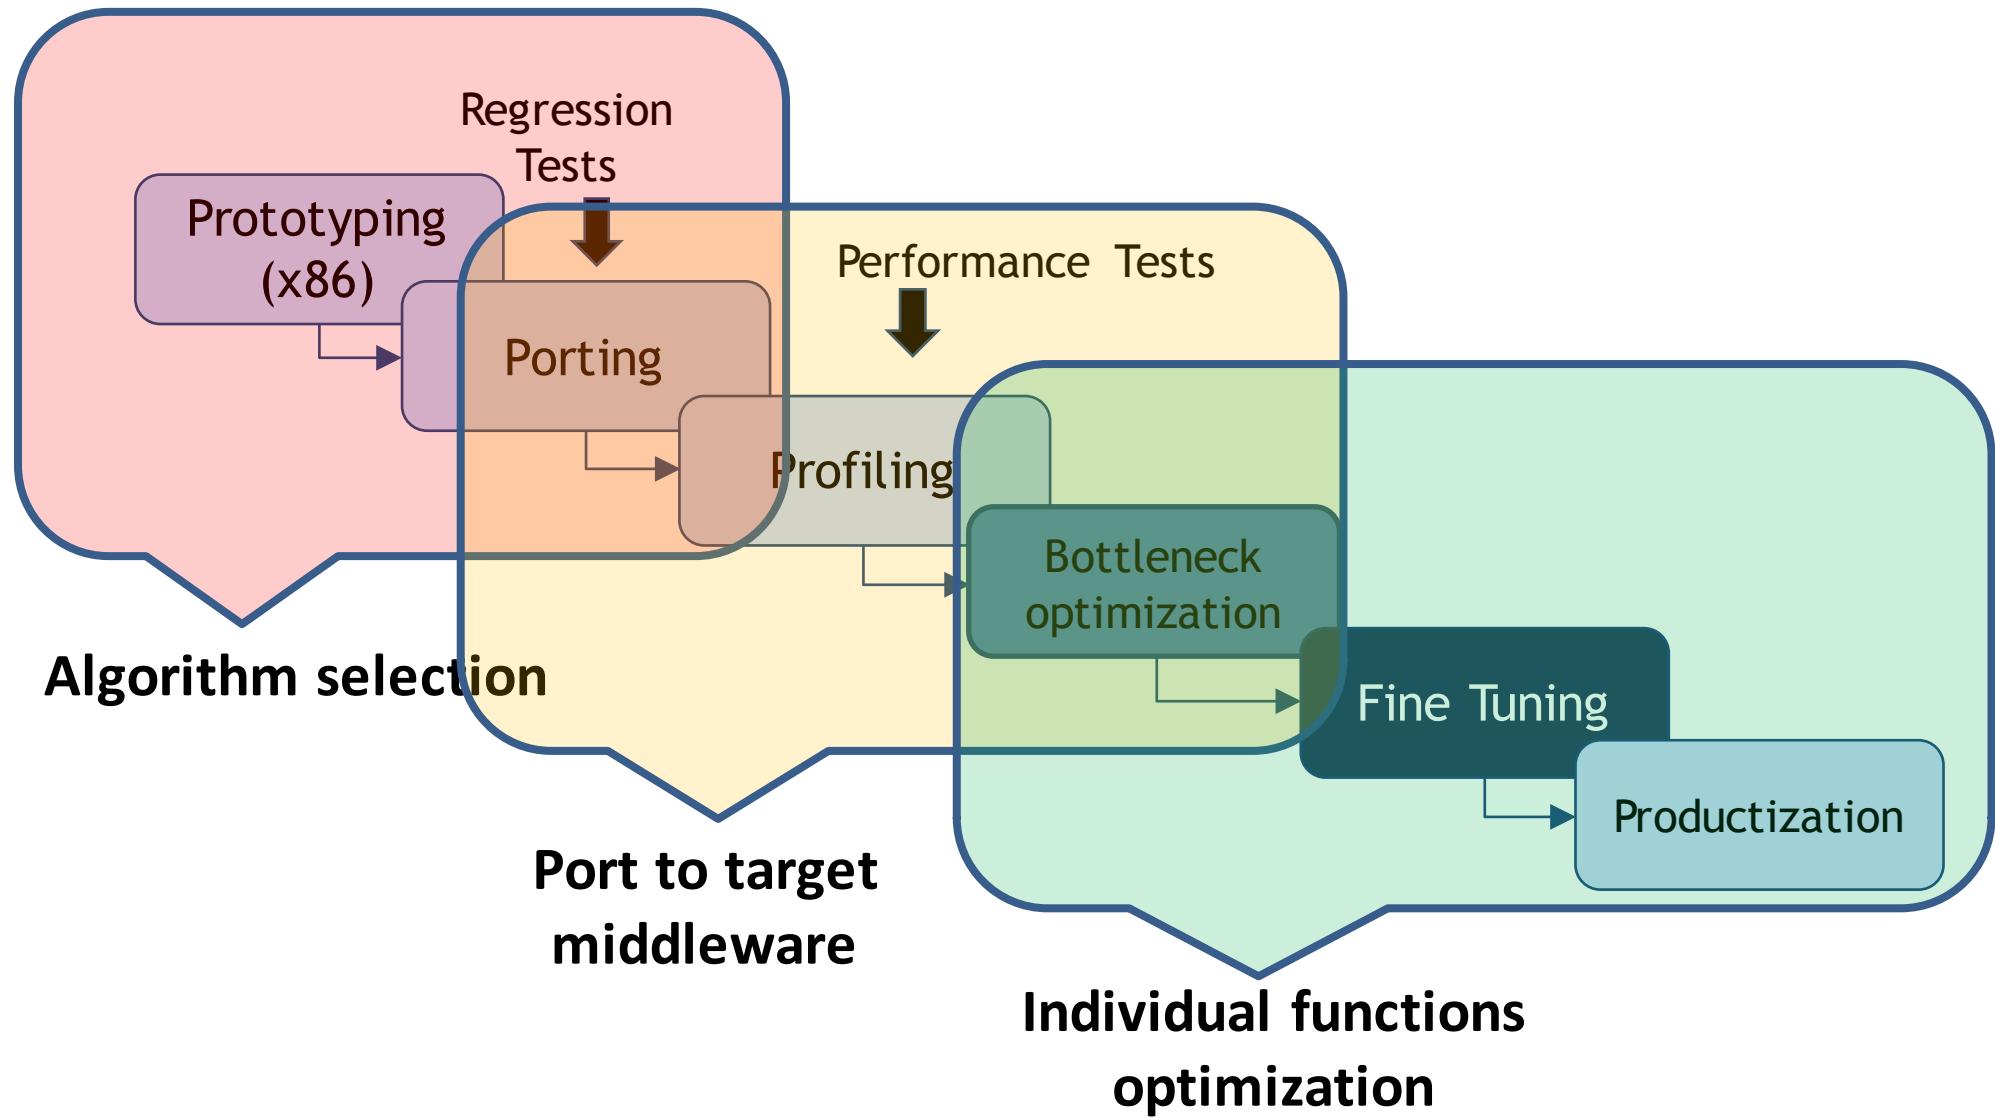


- Even cv::resize was too slow
- Wrappers to move images between android and opencv were not realtime

Embedded vision challenges

- Need to run in real-time on embedded/mobile/wearable devices
- Intense and power hungry computations
 - Heavy image processing requires significant memory bandwidth
 - Multiple cores do not help
- Mobile platforms are even more complicated
 - Unstable FPS impacts algorithm complexity (e.g. tracking)
 - Hard to predict when/if we are consuming too much
 - Thermal protection, power saving are hard to control and influence
- Very few computer vision accelerators

Embedded algorithm implementation

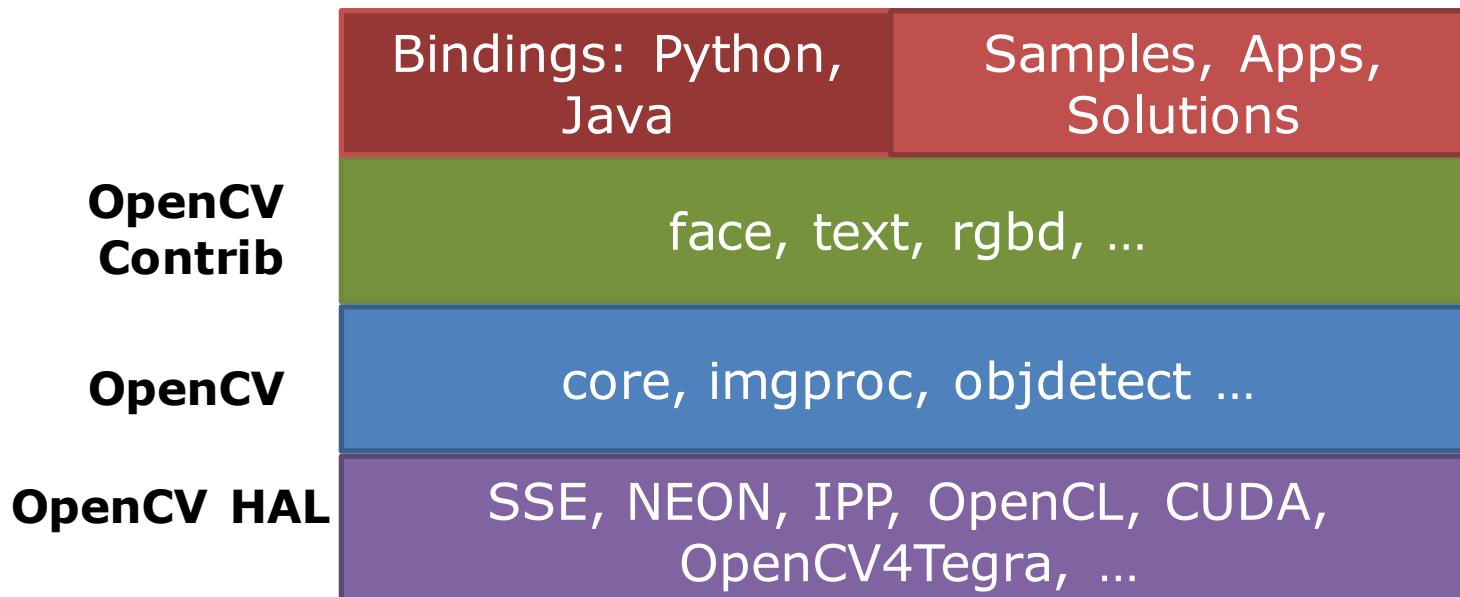


Available middleware frameworks

- Proprietary platforms (e.g. FastCV, VisionWorks)
- OpenCV
 - SSE, AVX
 - NEON
 - CUDA
 - OpenCL
 - Proprietary optimizations
- OpenVX

OpenCV at glance

- BSD license, **10M** downloads, **500K+** lines of code
- Huge community involvement, automated patch testing and integration process
- Runs everywhere



- Find more at <http://opencv.org> (user)
- Or <http://code.opencv.org> (developer)

OpenCV QA

Contribution/patch workflow:
see OpenCV wiki



<http://github.com/itseez/opencv>



<http://build.opencv.org>: buildbot with 50+ builders

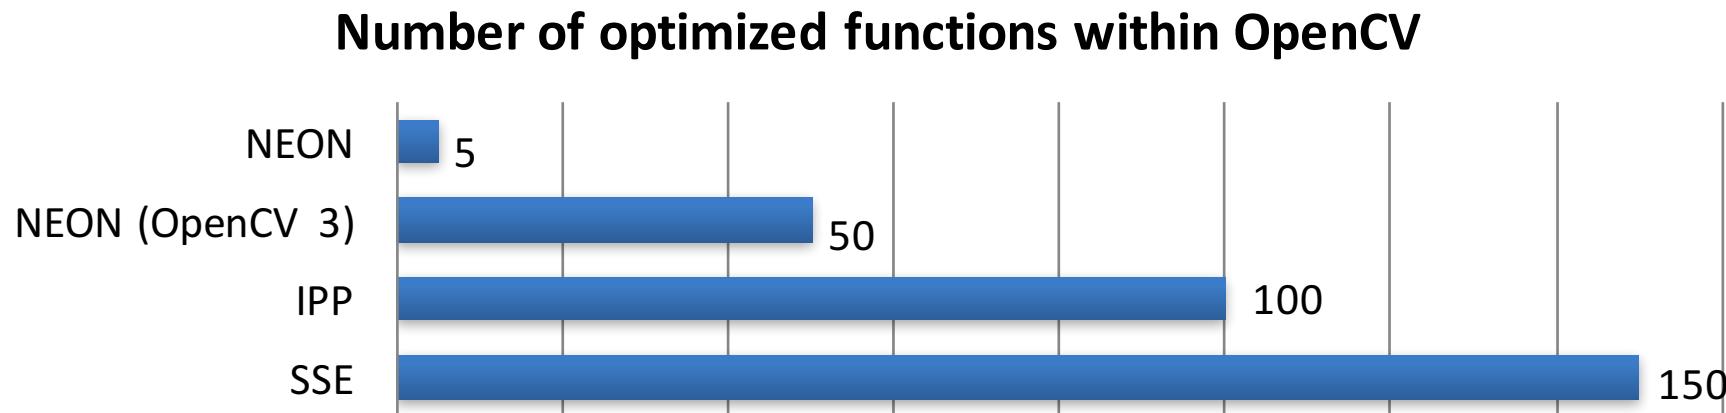


<http://pullrequest.opencv.org>: tests each patch

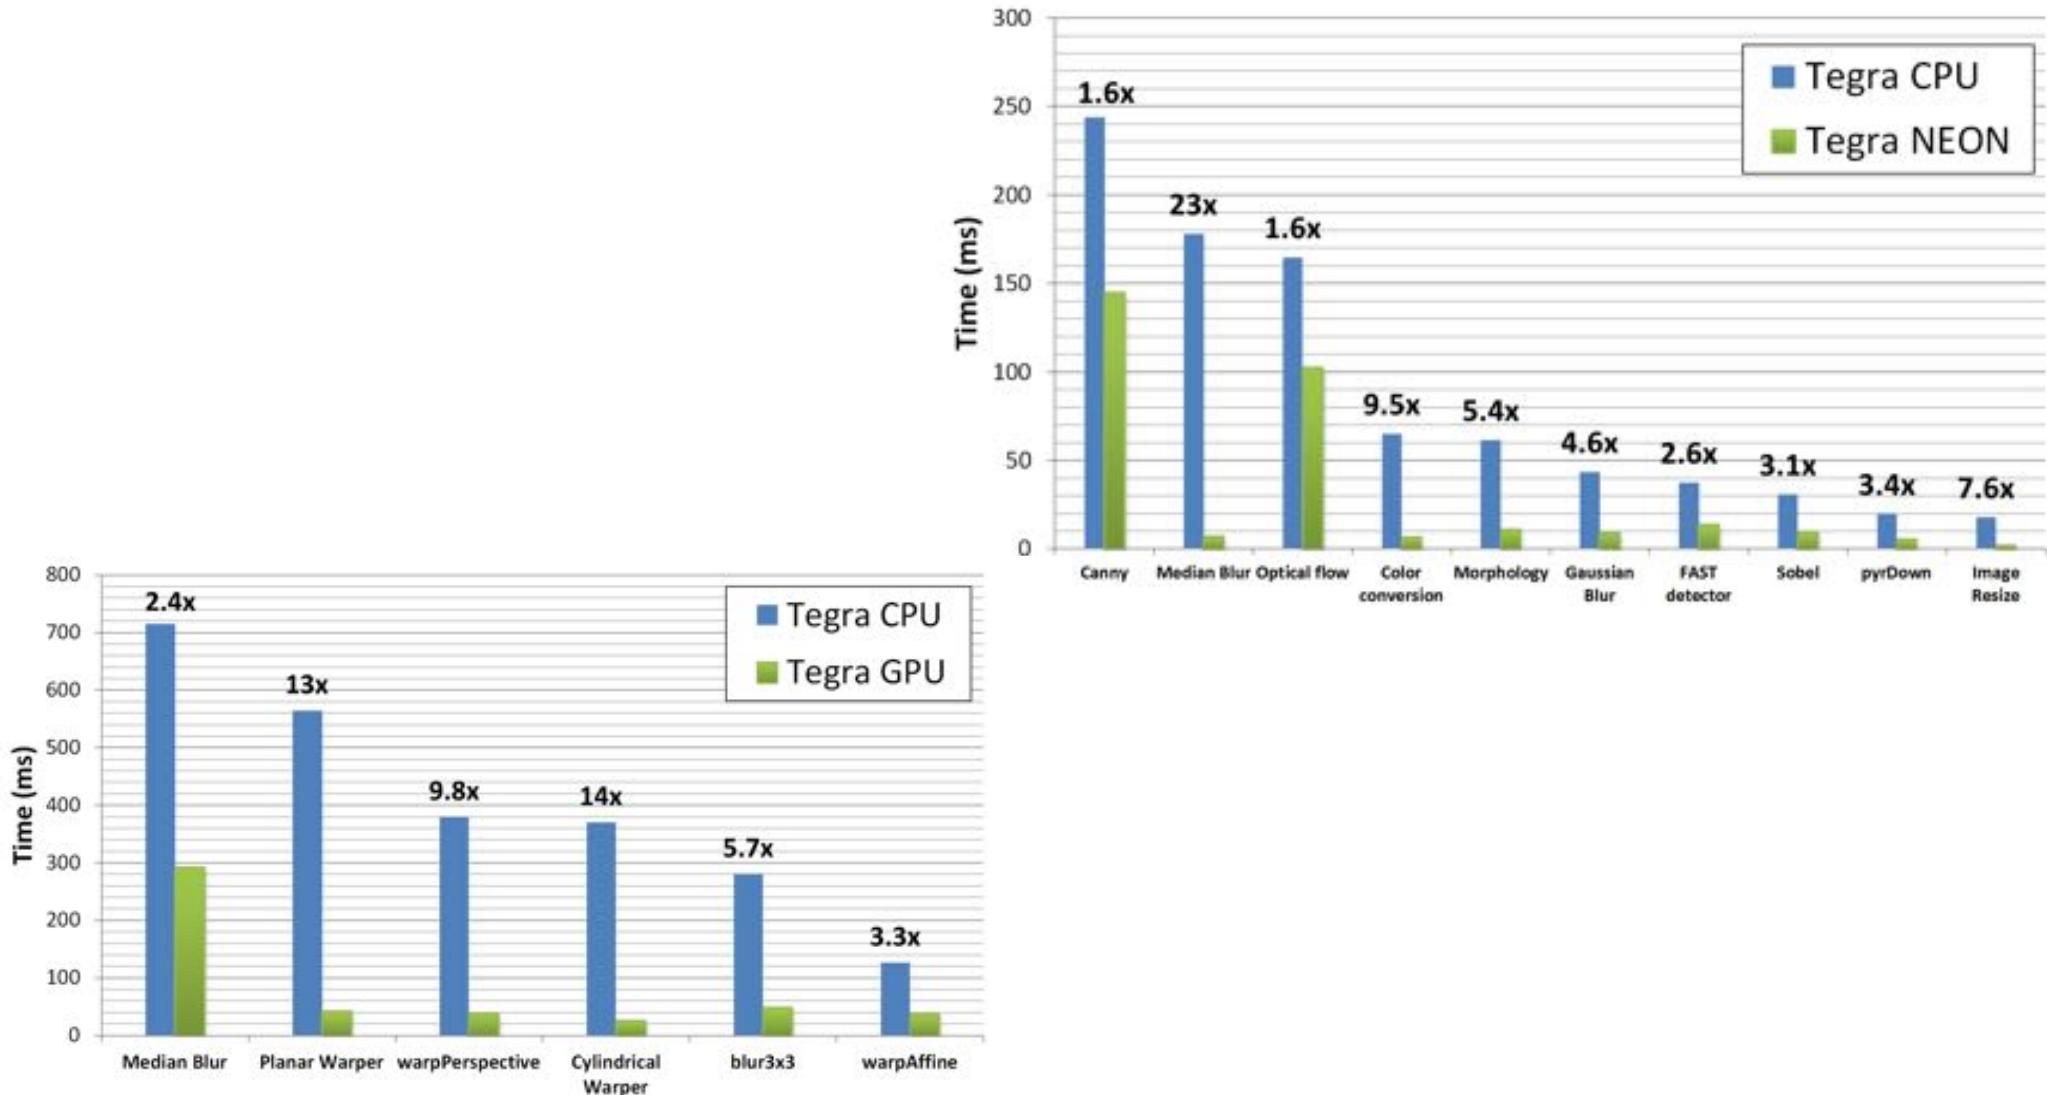


Why my OpenCV code is slow on mobile?

- OpenCV was initially optimized for desktop where it works fast
- ARM optimizations are far behind
- Scalar code does not perform on ARM as good as on x86
- Optimization might help to some extent

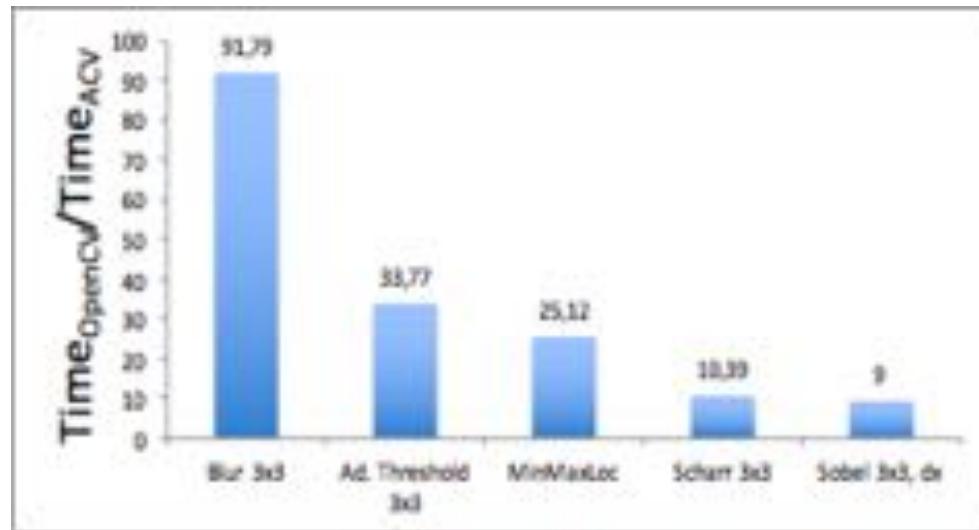


Early ARM/GLSL speedups*

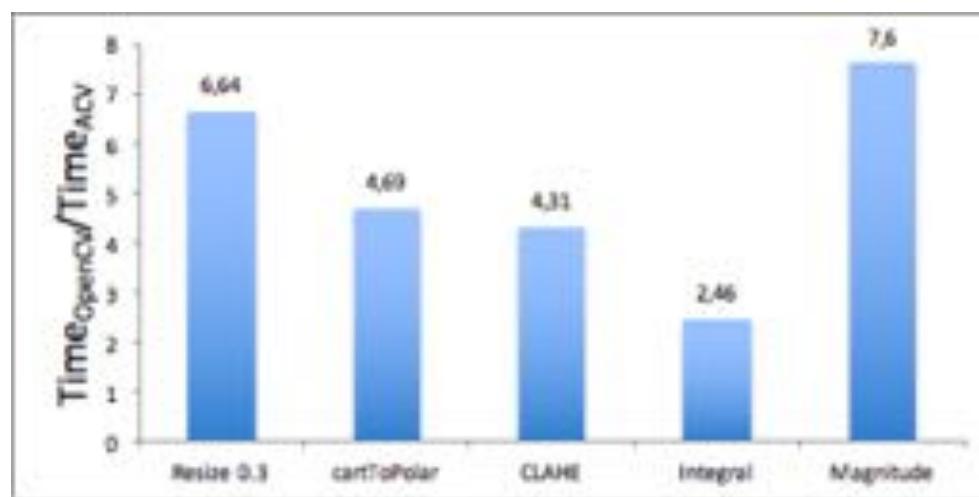


*Kari Puli et al, **Real-time computer vision with OpenCV**,
Communications of the ACM, Volume 55 Issue 6, June 2012, Pages 61-69

AcceleratedCV (Itseez) speedup against OpenCV



Comparison against public 2.4.9 release of OpenCV
(compiled with GCC 4.9, auto vectorization ON, -O3)
Qualcomm Snapdragon 800 SoC (2.15 Ghz, 4 cores), Android
Input image size is 720p



Higher level functions speedups

- **Robotics algorithms**
 - Acceleration up to **2x**
 - Resize, Filtering, CLAHE, blur, LBP cascade, Optical Flow, FAST features
- **ADAS algorithms**
 - Acceleration about **20-50% (1.25x-2x)**
 - Threshold, MinMaxLoc, Resize, Optical Flow, Match Template

What is wrong with OpenCV as middleware for embedded platforms?

- **Transparent memory model**
 - Users have to optimize memory throughput themselves
- **Does not address heterogeneous platform structure**
 - Users have to manage different accelerators
- **Many optimization opportunities are lost**
 - Choosing optimal accelerator
 - Filter stacking
 - Tiling



Vision Acceleration

Khronos Connects Software to Silicon

Open Consortium creating
ROYALTY-FREE, OPEN STANDARD
APIs for hardware acceleration

Defining the roadmap for
low-level silicon interfaces
needed on every platform

Graphics, compute, rich media,
vision, sensor and camera
processing

Rigorous specifications AND
conformance tests for cross-
vendor portability

*Acceleration APIs
BY the Industry
FOR the Industry*

<http://accelerateyourworld.org/>



Well over a *BILLION* people use Khronos APIs
Every Day...

Khronos Standards

K H R O N O S
G R O U P



Visual Computing

- 3D Graphics
- Heterogeneous Parallel Computing

3D Asset Handling

- 3D authoring asset interchange
- 3D asset transmission format with compression



Over 100 companies defining royalty-free APIs to connect software to silicon



Sensor Processing

- Vision Acceleration
- Camera Control
- Sensor Fusion



Acceleration in HTML5

- 3D in browser - no Plug-in
- Heterogeneous computing for JavaScript



We're a
Favorite Place
on Google



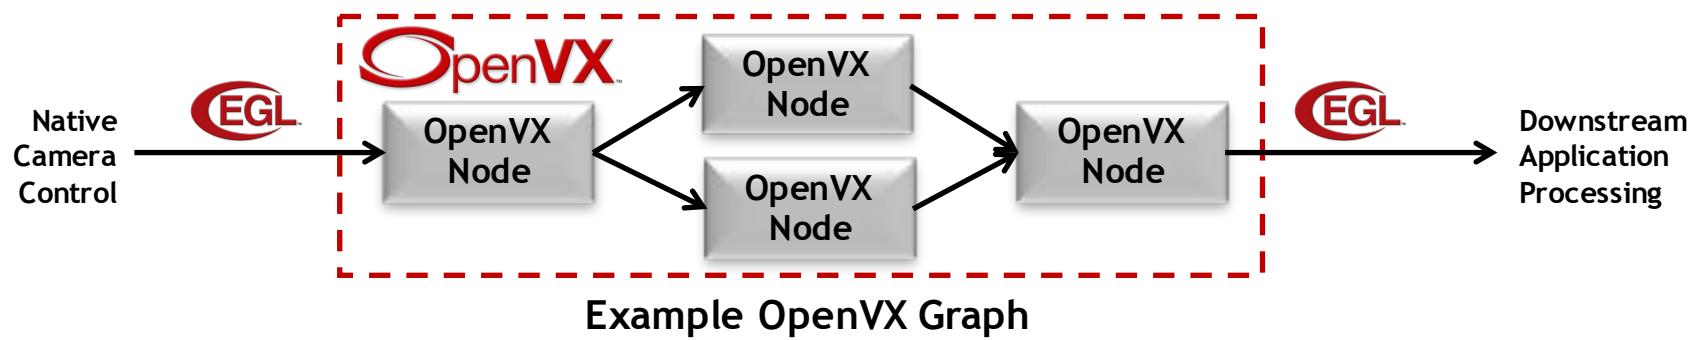
OpenVX - Power Efficient Vision Acceleration

- Out-of-the-Box vision acceleration framework
 - Enables low-power, real-time applications
 - Targeted at mobile and embedded platforms
- Functional Portability
 - Tightly defined specification
 - Full conformance tests
- Performance portability across diverse HW
 - Higher-level abstraction hides hardware details
 - ISPs, Dedicated hardware, DSPs and DSP arrays, GPUs, Multi-core CPUs ...
- Enables low-power, always-on acceleration
 - Can run solely on dedicated vision hardware
 - Does not require full SOC CPU/GPU complex to be powered on



OpenVX Graphs - The Key to Efficiency

- Vision processing directed graphs for power and performance efficiency
 - Each Node can be implemented in software or accelerated hardware
 - Nodes may be fused by the implementation to eliminate memory transfers
 - Processing can be tiled to keep data entirely in local memory/cache
- VXU Utility Library for access to single nodes
 - Easy way to start using OpenVX by calling each node independently
- Opaque memory model for images
 - Graph manager can move images between host memory and accelerator memory
 - Virtual images (users cannot see data)



OpenVX 1.0 Function Overview

- Core data structures
 - Images and Image Pyramids
 - Processing Graphs, Kernels, Parameters
- Image Processing
 - Arithmetic, Logical, and statistical operations
 - Multichannel Color and BitDepth Extraction and Conversion
 - 2D Filtering and Morphological operations
 - Image Resizing and Warping
- Core Computer Vision
 - Pyramid computation
 - Integral Image computation
- Feature Extraction and Tracking
 - Histogram Computation and Equalization
 - Canny Edge Detection
 - Harris and FAST Corner detection
 - Sparse Optical Flow

OpenVX Specification
Is Extensible

Khronos maintains extension registry

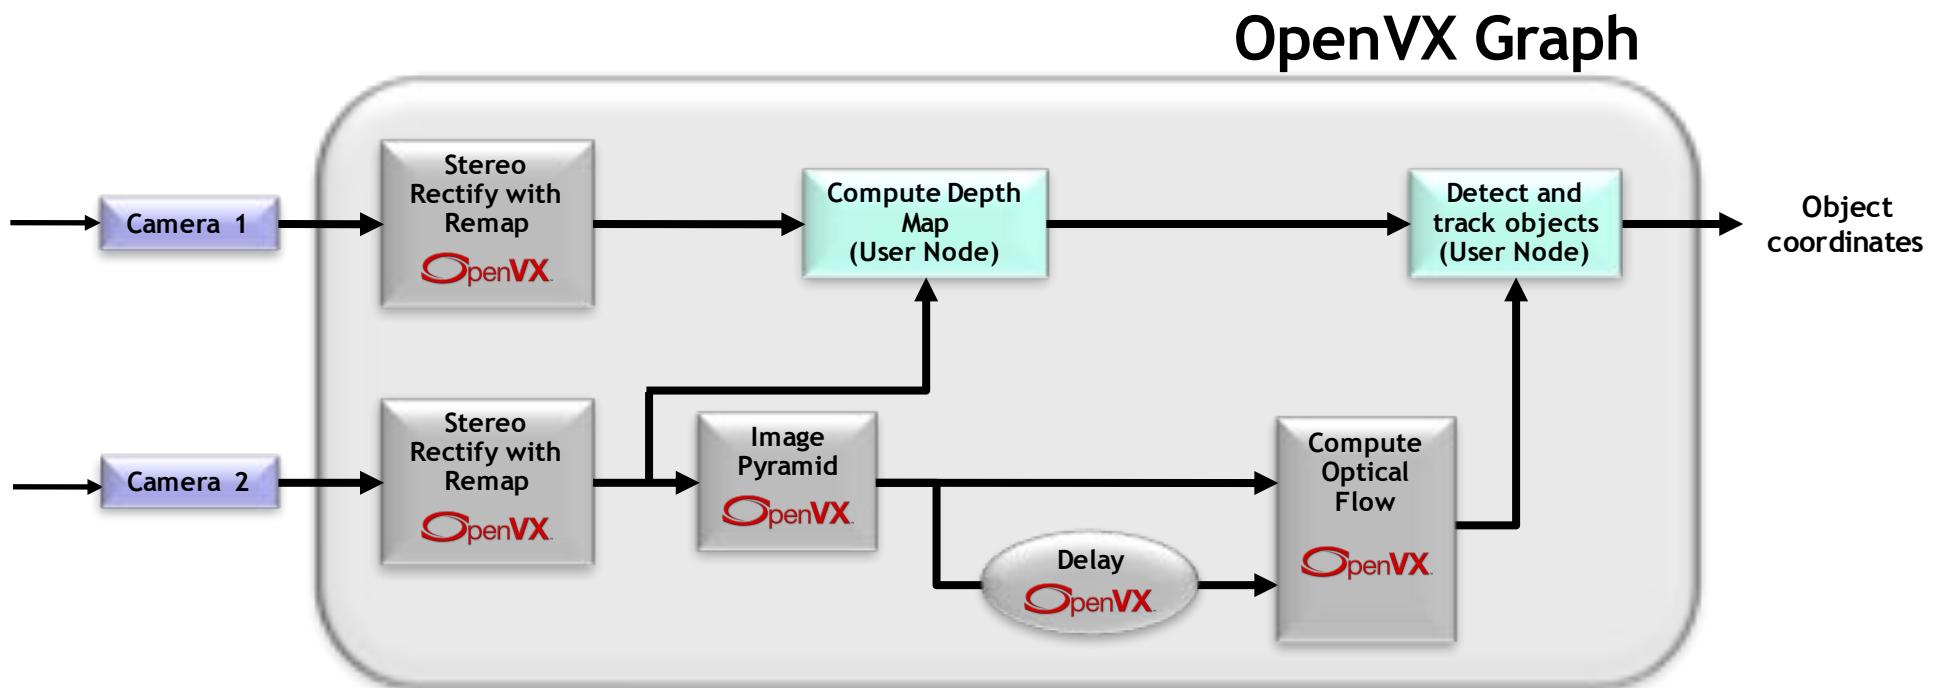
OpenVX 1.0 defines
framework for
creating, managing and
executing graphs

Focused set of widely
used functions that are
readily accelerated

Widely used extensions
adopted into future
versions of the core

Implementers can add
functions as extensions

Example Graph - Stereo Machine Vision



Tiling extension enables user nodes (extensions) to also optimally run in local memory

OpenVX standard

Specification document

- Defines the API
- Available to all in html and pdf from
<https://www.khronos.org/registry/vx/>
- A conformant implementation has to implement all the functions defined by the spec
- Extensions
 - Optional for implementation

Conformance tests

- Are used to enforce functional portability
- Available to adopters (\$20K Adopters fee, \$15K for members)
- Test suite exercises graph framework and functionality of each OpenVX 1.0 node
- OpenVX group reviews submitted results
- Approved Conformant implementations can use the OpenVX trademark

Sample implementation

- Open source implementation of OpenVX
- Available from
<https://www.khronos.org/registry/vx/>
- Conformant OpenVX implementation
- Was created to help create optimized OpenVX implementations

OpenVX

- OpenVX 1.1 specification released May 2016
 - www.khronos.org/openvx
- Full conformance test suite and Adopters Program immediately available
- Multiple companies have announced the release of OpenVX implementation, including Vivante, NVIDIA, Synopsys, Imagination, AMD
 - Many more to come!

amazon.com®

AMD

Aptina
IMAGING

ARM

AXIS®
COMMUNICATIONS

BDTⁱ

BROADCOM

cadence®

CEVA

cogniVUE®
CORPORATION

OMP

HUAWEI

intel

Imagination

itseez

KISHONI
INFORMATICS

Movidius

MULTICORE
WARE

NVIDIA.

QUALCOMM®

RENESAS

SAMSUNG

ST
life.augmented

TEXAS
INSTRUMENTS

videantis
passion for video

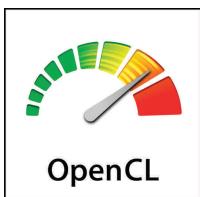
VIVANTE

XILINX®

OpenVX and OpenCV are Complementary

	 OpenCV	 OpenVX™
Governance	Community driven open source with no formal specification	Formal specification defined and implemented by hardware vendors
Conformance	No conformance tests for consistency and every vendor implements different subset	Full conformance test suite / process creates a reliable acceleration platform
Portability	APIs can vary depending on processor	Hardware abstracted for portability
Scope	Very wide 1000s of imaging and vision functions Multiple camera APIs/interfaces	Tight focus on hardware accelerated functions for mobile vision Use external camera API
Efficiency	Memory-based architecture Each operation reads and writes memory	Graph-based execution Optimizable computation, data transfer
Use Case	Rapid experimentation	Production development & deployment

Khronos APIs for Vision Processing



GPU Compute Shaders (OpenGL 4.X and OpenGL ES 3.1)

Pervasively available on almost any mobile device or OS

Easy integration into graphics apps - no vision/compute API interop needed

Program in GLSL not C

Limited to acceleration on a single GPU

General Purpose Heterogeneous Programming Framework

Flexible, low-level access to any devices with OpenCL compiler

Single programming and run-time framework for CPUs, GPUs, DSPs, hardware

Open standard for any device or OS - being used as backed by many languages and frameworks

Needs full compiler stack and IEEE precision



Out of the Box Vision Framework - Operators and graph framework library

Can run some or all modes on dedicated hardware - no compiler needed

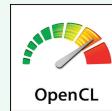
Higher-level abstraction means easier performance portability to diverse hardware

Graph optimization opens up possibility of low-power, always-on vision acceleration

Fixed set of operators - but can be extended

It is possible to use OpenCL or GLSL to build OpenVX Nodes on programmable devices!

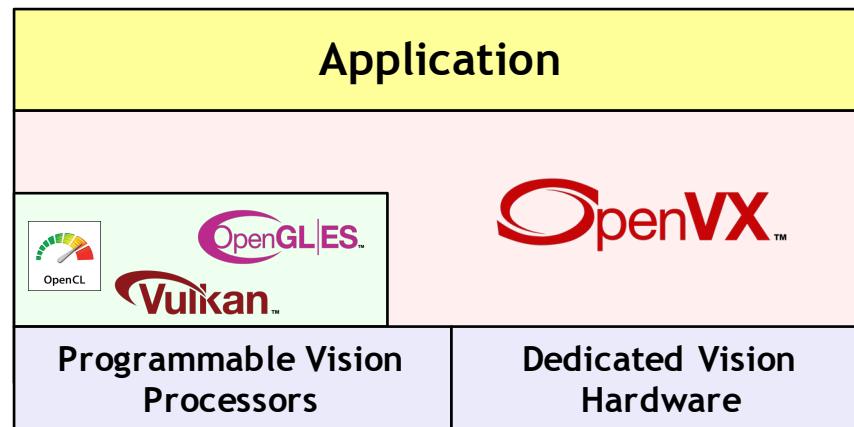
OpenVX and OpenCL are Complementary

		
Use Case	General Heterogeneous programming	Domain targeted vision processing
Ease of Use	General-purpose math libraries with no built-in vision functions	Fully implemented vision operators and framework 'out of the box'
Architecture	Language-based – needs online compilation	Library-based - no online compiler required
Target Hardware	'Exposed' architected memory model – can impact performance portability	Abstracted node and memory model - diverse implementations can be optimized for power and performance
Precision	Full IEEE floating point mandated	Requires minimal floating point or fixed point–optimized for vision operators

Layered Vision Processing Architecture

- The higher abstraction level of OpenVX protects app from hardware differences
 - Enables low-power, always-on acceleration - with application portability
- OpenVX is the Khronos vision API that does not NEED a CPU complex to be turned on during execution
 - Can use any processor for node execution - including hardware blocks
 - Can be configured to respond to incoming images with no host interaction

Implementers may choose to use OpenCL or OpenGL Compute Shaders to *implement* OpenVX nodes



And then use OpenVX to enable a developer to easily *connect* those nodes into a graph

OpenVX 1.0 extensions

- **Tiling extension: more efficient processing of graphs with user nodes**
 - Provisional spec released
- **XML Schema extension: cross-platform graph saving and loading**
 - Provisional spec released

Summary

- Khronos is building a trio of interoperating APIs for portable / power-efficient vision and sensor processing
- OpenVX 1.0.1 specification is now finalized and released
 - Full conformance tests and Adopters program immediately available
 - First commercial implementation released
- Any company is welcome to join Khronos to influence the direction of mobile and embedded vision processing!
 - \$15K annual membership fee for access to all Khronos API working groups
 - Well-defined IP framework protects your IP and conformant implementations
- More Information
 - www.khronos.org
 - ntrevett@nvidia.com
 - @neilt3d

Summary

- Exploding amount of computer vision applications running on mobile and embedded devices
- There is a need for an efficient software stack
 - OpenCV will help in the short term
 - OpenVX will simplify the problem in the long run
- Algorithm should be designed to work on embedded platforms from the very beginning
 - With functional safety requirements

Sample models



Traffic sign recognition: Existing approaches

- Color segmentation
- Grey-scale segmentation
- Shape fitting
- Viola-Jones cascade approach
- OCR
 - Text localization
 - Neumann, Matas, CVPR 2012

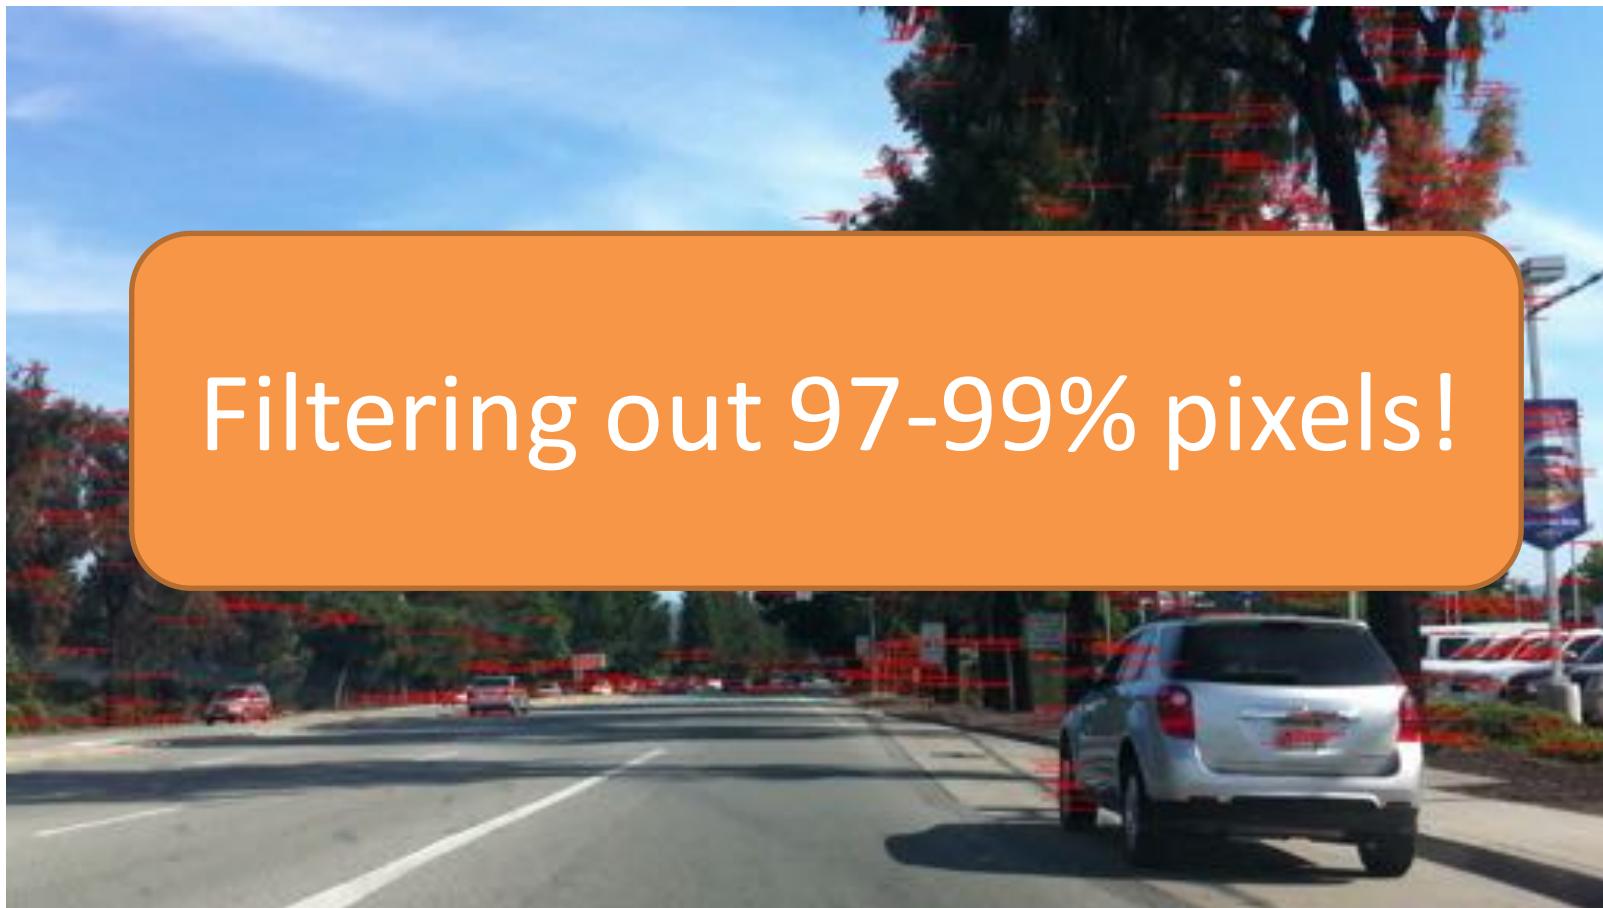
Can we detect a scanline that crosses a sign?



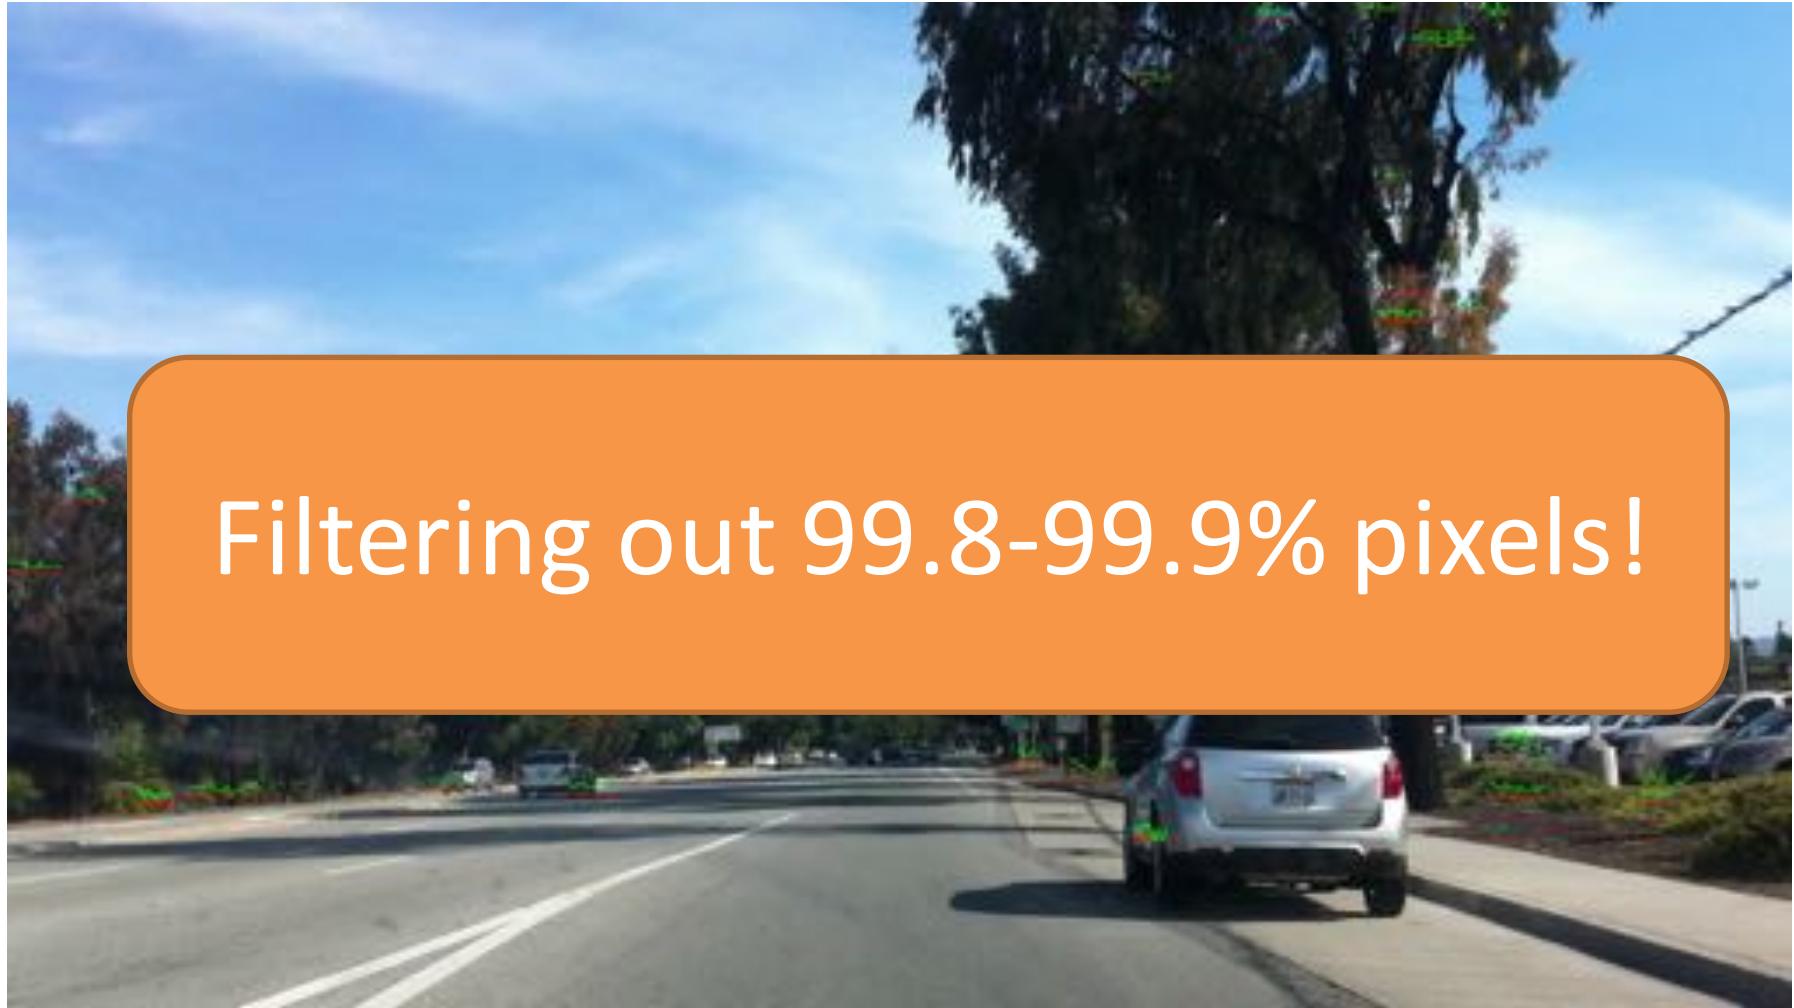
Algorithm: W filter

Detect a W sequence of extrema

Filtering out 97-99% pixels!



W + Classification filtering



Now, we have to analyze ROI



W + Classification filters out
~94% pixels!